

## Business Problem

Unlike most other reviewing websites which calculate the average of review scores (i.e. Amazon 4.5/5 stars, IMDB 8.2/10 rating), Rotten Tomatoes scores media on a percentage out of 100. This percentage is aggregated completely based on whether the review is positive or negative.

Rotten Tomatoes currently aggregates reviews in two ways based on:

1. Opinions of film and television critics
2. Opinions of users

Rotten Tomatoes wants to explore what the rest of the internet is saying about film and television. They want to have the ability to efficiently analyze publicly available comments/chatter (i.e. Tweets, Reddit message boards, Facebook comments) and derive a percentage score for movies and television programming.

This project will focus on just the first step of solving this business problem. Here is our proposed hypothesis:

'A model can be derived to input opinionated language material written about a film or show, so that it can reliably predict positive or negative sentiment.'

## Data Understanding

Our hypothesis will be tested on this [IMDB data set](https://www.kaggle.com/lakshmi25npathi/imdb-dataset-of-50k-movie-reviews) (<https://www.kaggle.com/lakshmi25npathi/imdb-dataset-of-50k-movie-reviews>) from kaggle. It contains 50,000 "highly polar" movie/show reviews and their positive or negative sentiment.

## Data Preparation

### Imports

```

In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import plotly.express as px

from bs4 import BeautifulSoup

import re

import nltk
from nltk.tokenize import word_tokenize
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer

from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
import string
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.pipeline import make_pipeline
from sklearn.model_selection import cross_val_score
from sklearn.metrics import accuracy_score
from sklearn.metrics import recall_score, precision_score, f1_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import plot_confusion_matrix
from sklearn.metrics import roc_auc_score, plot_roc_curve

```

```

In [2]: df = pd.read_csv('data/IMDB_Dataset.csv')

```

```

In [3]: df.head()

```

Out[3]:

	review	sentiment
0	One of the other reviewers has mentioned that ...	positive
1	A wonderful little production.   The...	positive
2	I thought this was a wonderful way to spend ti...	positive
3	Basically there's a family where a little boy ...	negative
4	Petter Mattei's "Love in the Time of Money" is...	positive

## Check for duplicates

```
In [4]: df.describe()
```

```
Out[4]:
```

	review	sentiment
count	50000	50000
unique	49582	2
top	Loved today's show!!! It was a variety and not...	positive
freq	5	25000

A small amount of duplicates. Since the data does not distinguish the media associated with the review, we will not remove any duplicates.

## Check Target Distribution

```
In [5]: df['sentiment'].value_counts()
```

```
Out[5]: positive    25000  
negative    25000  
Name: sentiment, dtype: int64
```

This is a balanced dataset

## Check Null Values

```
In [6]: df.isna().sum()
```

```
Out[6]: review    0  
sentiment    0  
dtype: int64
```

## Explore a single review

```
In [7]: df.loc[0][0]
```

```
Out[7]: "One of the other reviewers has mentioned that after watching just 1 Oz episode you'll be hooked. They are right, as this is exactly what happened with me.<br /><br />The first thing that struck me about Oz was its brutality and unflinching scenes of violence, which set in right from the word GO. Trust me, this is not a show for the faint hearted or timid. This show pulls no punches with regards to drugs, sex or violence. Its is hardcore, in the classic use of the word.<br /><br />It is called OZ as that is the nickname given to the Oswald Maximum Security State Penitentiary. It focuses mainly on Emerald City, an experimental section of the prison where all the cells have glass fronts and face inwards, so privacy is not high on the agenda. Emerald City is home to many..Aryans, Muslims, gangstas, Latinos, Christians, Italians, Irish and more....so scuffles, death stares, dodgy dealings and shady agreements are never far away.<br /><br />I would say the main appeal of the show is due to the fact that it goes where other shows wouldn't dare. Forget pretty pictures painted for mainstream audiences, forget charm, forget romance...OZ doesn't mess around. The first episode I ever saw struck me as so nasty it was surreal, I couldn't say I was ready for it, but as I watched more, I developed a taste for Oz, and got accustomed to the high levels of graphic violence. Not just violence, but injustice (crooked guards who'll be sold out for a nickel, inmates who'll kill on order and get away with it, well mannered, middle class inmates being turned into prison bitches due to their lack of street skills or prison experience) Watching Oz, you may become comfortable with what is uncomfortable viewing....thats if you can get in touch with your darker side."
```

## Text Cleaning

```
In [8]: def text_preprocessor(review):
# Strip html text
soup = BeautifulSoup(review, "lxml")
data = soup.get_text()
# Lower case all text
data = data.lower()
# Remove edge cases with multiple periods
pattern = re.compile(r'\s+')
data = re.sub(pattern, ' ', data.replace('.', ' '))
# Remove punctuation and other special characters
pattern = r'[^a-zA-Z\s]'
data = re.sub(pattern, '', data)
# Tokenize
data = word_tokenize(data)
# Get list of nltk's stopwords
stopwords_list = stopwords.words('english')
# Remove stop words
data = [word for word in data if word not in stopwords_list]
# Initialize a PorterStemmer object
stemmer = PorterStemmer()
# Convert the tokens into their stem
data = [stemmer.stem(token) for token in data]
# Convert the list of words back into
# a string by joining each word with a space
data = ' '.join(data)
# Remove double spaces
data = data.replace(' ', ' ')
# Remove opening and trailing spaces
data = data.strip()
# Return the cleaned text data
return data
```

```
In [9]: df.review = df.review.apply(text_preprocessor)
```

```
In [10]: print(df.review[:5])
```

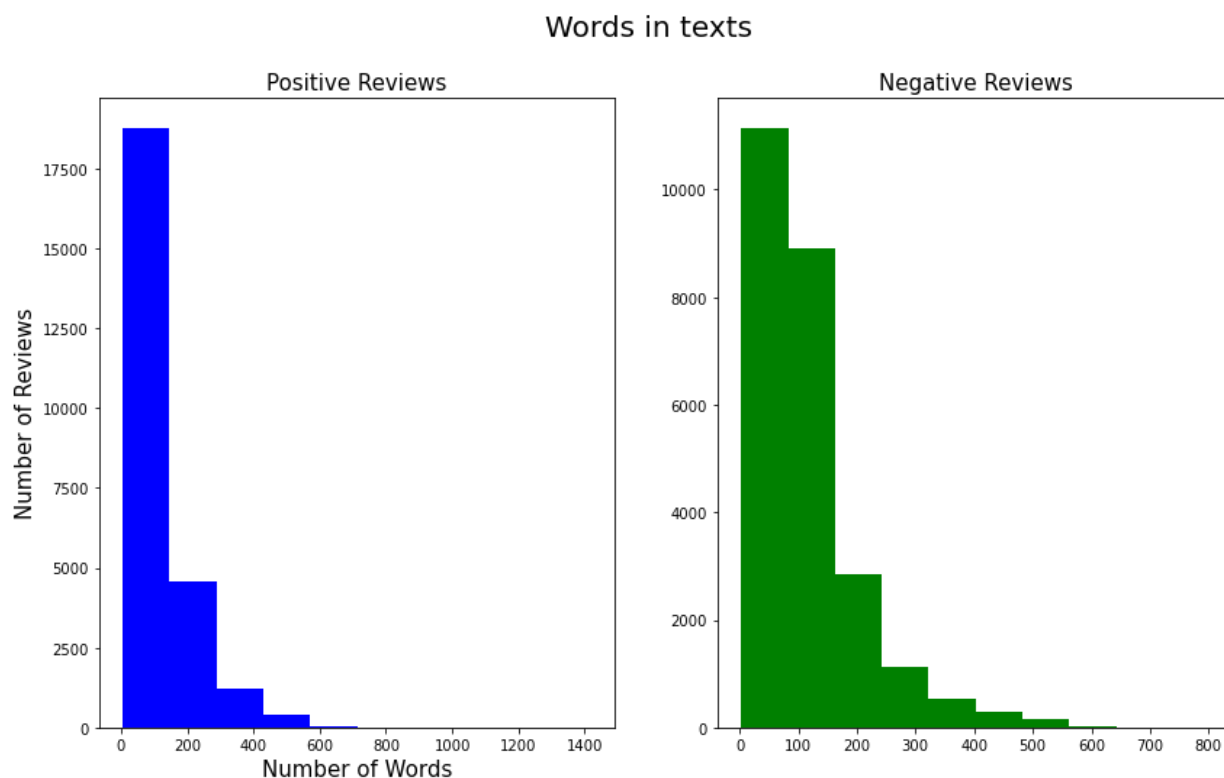
```
0    one review mention watch oz episod youll hook ...
1    wonder littl product film techniqu unassum old...
2    thought wonder way spend time hot summer weeke...
3    basic there famili littl boy jake think there ...
4    petter mattei love time money visual stun film...
Name: review, dtype: object
```

## Data Exploration

Now that we have cleaned our text, we will analyze some of our data's characteristics:

### Number of words per review

```
In [11]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(14, 8))
length_good_reviews = df[df['sentiment'] == 'positive']['review'].str.split()
ax1.hist(length_good_reviews, color='blue')
ax1.set_title('Positive Reviews', fontsize=15)
ax1.set_ylabel('Number of Reviews', fontsize=15)
ax1.set_xlabel('Number of Words', fontsize=15)
length_bad_reviews = df[df['sentiment'] == 'negative']['review'].str.split()
ax2.hist(length_bad_reviews, color='green')
ax2.set_title('Negative Reviews', fontsize=15)
fig.suptitle('Words in texts', fontsize=20)
plt.show()
```



It seems that negative reviews are generally more verbose than positive reviews.

### Mean word length per review

```
In [12]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(22, 8))
pos_word = df[df['sentiment'] == 'positive']['review'].str.split().apply(la
len(i) for i in x])
sns.distplot(pos_word.map(lambda x: np.mean(x)), ax=ax1, color='blue')
ax1.set_title('Positive Reviews')
neg_word = df[df['sentiment'] == 'negative']['review'].str.split().apply(la
len(i) for i in x])
sns.distplot(neg_word.map(lambda x: np.mean(x)), ax=ax2, color='green')
ax2.set_title('Negative Reviews')
fig.suptitle('Mean word length per review', fontsize=20)
```

/Users/dan/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages/seaborn/distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


warnings.warn(msg, FutureWarning)

/Users/dan/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages/seaborn/distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

Out[12]: Text(0.5, 0.98, 'Mean word length per review')

Mean word length per review



Word length in positive and negative reviews are relatively equal.

For further exploration, we must retrieve a corpus of all words:

## Grab all words in reviews

```
In [13]: def get_all_words(text):
          words = []
          for x in text:
              for y in x.split():
                  words.append(y.strip())
          return words
all_words = get_all_words(df.review)
all_words[:10]
```

```
Out[13]: ['one',
          'review',
          'mention',
          'watch',
          'oz',
          'episod',
          'youll',
          'hook',
          'right',
          'exactli']
```

## Count all words in reviews

```
In [14]: from collections import Counter
          counter = Counter(all_words)
          most_common = counter.most_common(10)
          most_common = dict(most_common)
          most_common
```

```
Out[14]: {'movi': 101385,
          'film': 94157,
          'one': 53941,
          'like': 44144,
          'time': 30810,
          'good': 29423,
          'make': 28675,
          'character': 28038,
          'see': 27867,
          'get': 27820}
```

## Create a function to find most common n-grams

credit: <https://www.kaggle.com/madz2000/sentiment-analysis-cleaning-eda-bert-88-acc/notebook>  
(<https://www.kaggle.com/madz2000/sentiment-analysis-cleaning-eda-bert-88-acc/notebook>)




```
In [15]: def get_top_text_ngrams(text, num_words, ngram):
vec = TfidfVectorizer(ngram_range=(ngram, ngram)).fit(text)
bag_of_words = vec.transform(text)
sum_words = bag_of_words.sum(axis=0)
words_freq = [(word, sum_words[0, idx])
               for word, idx in vec.vocabulary_.items()]
words_freq = sorted(words_freq, key=lambda x: x[1], reverse=True)
return words_freq[:num_words]
```

## Unigram Analysis


```
In [16]: most_common_uni_pos = get_top_text_ngrams(df.review[df.sentiment=='positive'])
most_common_uni_pos = dict(most_common_uni_pos)
temp = pd.DataFrame(columns=["Weighted_words", 'TFIDF_Score'])
temp["Weighted_words"] = list(most_common_uni_pos.keys())
temp["TFIDF_Score"] = list(most_common_uni_pos.values())
fig = px.bar(temp, x="TFIDF_Score", y="Weighted_words", title='Weighted Words in Positive Reviews',
              width=700, height=700, color='Weighted_words')
fig.show()
```

## Weighted Words in Positive Reviews



```
In [17]: most_common_uni_neg = get_top_text_ngrams(df.review[df.sentiment=='negative'])
most_common_uni_neg = dict(most_common_uni_neg)
temp = pd.DataFrame(columns=["Weighted_words", 'TFIDF_Score'])
temp["Weighted_words"] = list(most_common_uni_neg.keys())
temp["TFIDF_Score"] = list(most_common_uni_neg.values())
fig = px.bar(temp, x="TFIDF_Score", y="Weighted_words", title='Weighted Words',
              width=700, height=700, color='Weighted_words')
fig.show()
```

## Weighted Words in Negative Reviews




## Bigram Analysis

```
In [18]: most_common_bi_pos = get_top_text_ngrams(df.review[df.sentiment=='positive'])
most_common_bi_pos = dict(most_common_bi_pos)
temp = pd.DataFrame(columns = ["Weighted_words" , 'TFIDF_Score'])
temp["Weighted_words"] = list(most_common_bi_pos.keys())
temp["TFIDF_Score"] = list(most_common_bi_pos.values())
fig = px.bar(temp, x="TFIDF_Score", y="Weighted_words", title='Weighted Bigrams in Positive Reviews',
width=700, height=700,color='Weighted_words')
fig.show()
```




## Weighted Bigrams in Positive Reviews



```
In [19]: most_common_bi_neg = get_top_text_ngrams(df.review[df.sentiment=='negative']
most_common_bi_neg = dict(most_common_bi_neg)
temp = pd.DataFrame(columns = ["Weighted_words" , 'TFIDF_Score'])
temp["Weighted_words"] = list(most_common_bi_neg.keys())
temp["TFIDF_Score"] = list(most_common_bi_neg.values())
fig = px.bar(temp, x="TFIDF_Score", y="Weighted_words", title='Weighted Bigrams in Negative Reviews',
width=700, height=700,color='Weighted_words')
fig.show()
```


## Weighted Bigrams in Negative Reviews



## Trigram Analysis


```
In [20]: most_common_tri_pos = get_top_text_ngrams(df.review[df.sentiment=='positive'])
most_common_tri_pos = dict(most_common_tri_pos)
temp = pd.DataFrame(columns = ["Weighted_words" , 'TFIDF_Score'])
temp["Weighted_words"] = list(most_common_tri_pos.keys())
temp["TFIDF_Score"] = list(most_common_tri_pos.values())
fig = px.bar(temp, x="TFIDF_Score", y="Weighted_words", title='Weighted Tri
width=700, height=700,color='Weighted_words')
fig.show()
```

## Weighted Trigrams Positive Reviews



```
In [21]: most_common_tri_neg = get_top_text_ngrams(df.review[df.sentiment=='negative'])
most_common_tri_neg = dict(most_common_tri_neg)
temp = pd.DataFrame(columns = ["Weighted_words" , 'TFIDF_Score'])
temp["Weighted_words"] = list(most_common_tri_neg.keys())
temp["TFIDF_Score"] = list(most_common_tri_neg.values())
fig = px.bar(temp, x="TFIDF_Score", y="Weighted_words", title='Weighted Tri
width=700, height=700,color='Weighted_words')
fig.show()
```

## Weighted Trigrams in Negative Reviews



## Modeling

## Model-less Baseline

```
In [22]: df['sentiment'].value_counts()
```

```
Out[22]: positive    25000  
negative    25000  
Name: sentiment, dtype: int64
```

Our model-less baseline says we can accurately predict the each class 50% success rate if we only guess one class every time.

## Encode Target Column

```
In [23]: df.sentiment.replace("positive" , 1 , inplace = True)  
df.sentiment.replace("negative" , 0 , inplace = True)
```

## Train test split

```
In [24]: X_train, X_test, y_train, y_test = train_test_split(df[['review']], df.sentiment)
```

## Initialize count and tfidf vectorizers

```
In [25]: count = CountVectorizer()  
tfidf = TfidfVectorizer(max_features=12500)
```

## Create pipelines containing each vectorizer

```
In [26]: models = {'lr_count': make_pipeline(count, LogisticRegression(max_iter=375,  
    'dt_count': make_pipeline(count, DecisionTreeClassifier(random_st  
    'rf_count': make_pipeline(count, RandomForestClassifier(random_st  
    'lr_tfidf': make_pipeline(tfidf, LogisticRegression(max_iter=375,  
    'dt_tfidf': make_pipeline(tfidf, DecisionTreeClassifier(random_st  
    'rf_tfidf': make_pipeline(tfidf, RandomForestClassifier(random_st
```

```
In [59]: models = {'lr_tfidf': make_pipeline(tfidf, LogisticRegression(max_iter=375,
```

We will prioritize F1 score in model selection for well balanced correctness



```
In [27]: baseline_scores = {}

for model in models:
    score = cross_val_score(models[model], X_train.iloc[:,0], y_train, scoring='f1')
    baseline_scores[model] = score.mean()

baseline_scores
```

```
Out[27]: {'lr_count': 0.8756128986644571,
          'dt_count': 0.7171327201737887,
          'rf_count': 0.8513636098212689,
          'lr_tfidf': 0.8900022620962762,
          'dt_tfidf': 0.7106417503189231,
          'rf_tfidf': 0.8460468712274147}
```

## Fit all training data on logistic regression model with tfidf vectorization

```
In [28]: lr_pipeline = models['lr_tfidf']
```

```
In [29]: lr_pipeline.fit(X_train.iloc[:,0], y_train)
```

```
Out[29]: Pipeline(steps=[('tfidfvectorizer', TfidfVectorizer(max_features=12500)),
                          ('logisticregression',
                           LogisticRegression(max_iter=375, random_state=42))])
```

## Evaluate the model

```
In [30]: def evaluate(estimator, X_train, X_test, y_train, y_test, roc_auc='proba'):
    """
    Evaluation function to show a few scores for both the train and test set
    Also shows a confusion matrix for the test set

    roc_auc allows you to set how to calculate the roc_auc score:
    'dec' for decision_function or 'proba' for predict_proba
    If roc_auc == 'skip', then it ignores calculating the roc_auc_score

    Function takes in:
    'estimator' a fit sklearn model object
    'X_train' dataframe
    'X_test' dataframe
    'y_train' series
    'y_test' series
    'roc_auc' string that defines how the score is calculated
    """

    # grab predictions
    train_preds = estimator.predict(X_train)
    test_preds = estimator.predict(X_test)

    # output needed for roc_auc_score
    if roc_auc == 'skip': # skips calculating the roc_auc_score
        train_out = False
        test_out = False
    elif roc_auc == 'dec': # not all classifiers have decision_function
        train_out = estimator.decision_function(X_train)
        test_out = estimator.decision_function(X_test)
    elif roc_auc == 'proba':
        train_out = estimator.predict_proba(
            X_train)[: , 1] # proba for the 1 class
        test_out = estimator.predict_proba(X_test)[: , 1]
    else:
        raise Exception(
            "The value for roc_auc should be 'skip', 'dec' or 'proba'.")

    # print scores
    print("Train Scores")
    print("-----")
    print(f"Accuracy: {accuracy_score(y_train, train_preds)}")
    print(f"Precision: {precision_score(y_train, train_preds)}")
    print(f"Recall: {recall_score(y_train, train_preds)}")
    print(f"F1 Score: {f1_score(y_train, train_preds)}")
    if type(train_out) == np.ndarray: # checking for roc_auc
        print(f"ROC-AUC: {roc_auc_score(y_train, train_out)}")
    print("----" * 5)
    print("Test Scores")
    print("-----")
    print(f"Accuracy: {accuracy_score(y_test, test_preds)}")
    print(f"Precision: {precision_score(y_test, test_preds)}")
    print(f"Recall: {recall_score(y_test, test_preds)}")
    print(f"F1 Score: {f1_score(y_test, test_preds)}")
    if type(test_out) == np.ndarray:
        print(f"ROC-AUC: {roc_auc_score(y_test, test_out)}")

    # plot test confusion matrix
```

```
plot_confusion_matrix(estimator, X_test, y_test, values_format=',.5g')
plt.grid(False)
plt.show()
```


```
In [31]: evaluate(lr_pipeline, X_train.iloc[:,0], X_test.iloc[:,0], y_train, y_test,
```

Train Scores

```
-----
Accuracy: 0.9219733333333333
Precision: 0.9143444134225359
Recall: 0.9303210591199014
F1 Score: 0.9222635494155155
ROC-AUC: 0.9754639240177175
-----
```

Test Scores

```
-----
Accuracy: 0.89072
Precision: 0.8844430712189093
Recall: 0.9025697619422985
F1 Score: 0.893414481897628
ROC-AUC: 0.9592380531179882
-----
```



The model is slightly overfit, but overall performed well on all metric scores. Further optimizations can be made on:

1. Vectorizer 'max features' hyperparameter
2. Inverse of regularization strength 'C' hyperparameter

## Grab Feature Importances

```
In [32]: coefs = lr_pipeline.steps[1][1].coef_
```

## Transform test data with fit tfidf vectorizer

```
In [33]: # the fit tfidf vectorizer  
transformer = lr_pipeline.steps[0][-1]
```

```
In [34]: X_transformed = transformer.transform(X_test.iloc[:,0]).toarray()
```

## Analyze coefficients for feature importance

### Words in negative reviews

```
In [63]: # Zip the names of the features with the features importance  
coef_magnitudes_neg = zip(transformer.get_feature_names(), coefs.squeeze()).
```

```
In [64]: # Sort the features in descending order by magnitude
top_100_neg = sorted(coef_magnitudes_neg, key=lambda x: x[1], reverse=False)
top_100_neg
```

```
Out[64]: [('worst', -9.979429095516444),
 ('wast', -8.897572147775568),
 ('aw', -7.69853494545529),
 ('bad', -7.499133681247096),
 ('bore', -6.970010664565677),
 ('poor', -6.241712171490241),
 ('terribl', -5.836886145433135),
 ('noth', -5.700127568440428),
 ('disappoint', -5.662935050120046),
 ('fail', -5.333368612472158),
 ('wors', -5.003899037007581),
 ('horribl', -4.909406708383349),
 ('poorli', -4.805561215512351),
 ('dull', -4.663559073580145),
 ('suppos', -4.587736246090588),
 ('lack', -4.520590601432766),
 ('save', -4.369268990324432),
 ('minut', -4.25737737396317),
 ('ridicul', -4.154887819324281),
 ('annoy', -4.1489427378412955),
 ('unfortun', -4.08775159102494),
 ('stupid', -3.9511499799907797),
 ('instead', -3.9476888858738),
 ('script', -3.8761877053583014),
 ('lame', -3.8244380780766156),
 ('mediocr', -3.59748632075463),
 ('mess', -3.5907856582381092),
 ('oh', -3.463797552047383),
 ('sorri', -3.4013464970450125),
 ('attempt', -3.374591976446817),
 ('avoid', -3.366955694201297),
 ('unless', -3.3534971942596212),
 ('would', -3.352055300496341),
 ('even', -3.3071164884309168),
 ('pointless', -3.2872214235273014),
 ('laughabl', -3.275749975751101),
 ('embarrass', -3.2472211299561224),
 ('mstk', -3.2221473139712034),
 ('unfunni', -3.1689461333045124),
 ('pathet', -3.1635056839400915),
 ('forgett', -3.099705735731295),
 ('badli', -3.0794857972828695),
 ('couldnt', -3.058838375942007),
 ('clich', -3.041269085640917),
 ('weak', -2.9821310573662867),
 ('redeem', -2.9752543895109658),
 ('tri', -2.955249073424983),
 ('reason', -2.84835483761114),
 ('least', -2.824204264922108),
 ('money', -2.7921741545090266),
 ('cheap', -2.790914505548369),
 ('predict', -2.7530521488818906),
 ('pretenti', -2.728398775779035),
```

```
(
    'idea', -2.721485270466022),
    ('neither', -2.717534644293169),
    ('look', -2.704203244427797),
    ('insult', -2.586591666586116),
    ('director', -2.5780439001501687),
    ('bother', -2.5576400708520146),
    ('stereotyp', -2.5435016446263132),
    ('tediou', -2.51282670236836),
    ('lousi', -2.510557932691127),
    ('suck', -2.5088161434253764),
    ('could', -2.5084057834982434),
    ('plot', -2.494387325283256),
    ('uninterest', -2.4794902451081144),
    ('none', -2.4653383319337916),
    ('garbag', -2.4601921530295576),
    ('wooden', -2.458294588474702),
    ('mildli', -2.4511246173785466),
    ('crap', -2.426146898260964),
    ('dread', -2.406355012024241),
    ('turkey', -2.3940695260134404),
    ('unintent', -2.379997239726048),
    ('far', -2.3770522403760572),
    ('seem', -2.37049640505965),
    ('potenti', -2.3633426716516737),
    ('silli', -2.316531737463489),
    ('miscast', -2.311969201564295),
    ('flat', -2.3030399678730293),
    ('problem', -2.2655435802052653),
    ('bland', -2.2425715312295553),
    ('enough', -2.236322351586425),
    ('pain', -2.192616361628102),
    ('materi', -2.190796043145468),
    ('stinker', -2.184538444692382),
    ('hour', -2.183738560768064),
    ('shallow', -2.171984850818519),
    ('nonsens', -2.165151288596038),
    ('seagal', -2.1449484825305825),
    ('uninspir', -2.139798404759646),
    ('irrit', -2.1378972603014796),
    ('effort', -2.136816009953017),
    ('amateurish', -2.131094478525395),
    ('dumb', -2.112446945273682),
    ('half', -2.0956442183228545),
    ('unbeliev', -2.0669049305979774),
    ('sit', -2.0597838653986424),
    ('write', -2.059154970182335),
    ('skip', -2.0486711796827706)]
```


```
In [68]: coef_magnitudes_neg = zip(transformer.get_feature_names(), coefs.squeeze().
neg_sorted = sorted(coef_magnitudes_neg, key=lambda x: x[1], reverse=False)
```

```
In [69]: neg_sorted
```

```
Out[69]: [('worst', -9.979429095516444),  
          ('wast', -8.897572147775568),  
          ('aw', -7.69853494545529),  
          ('bad', -7.499133681247096),  
          ('bore', -6.970010664565677),  
          ('poor', -6.241712171490241),  
          ('terribl', -5.836886145433135),  
          ('noth', -5.700127568440428),  
          ('disappoint', -5.662935050120046),  
          ('fail', -5.333368612472158),  
          ('wors', -5.003899037007581),  
          ('horribl', -4.909406708383349),  
          ('poorli', -4.805561215512351),  
          ('dull', -4.663559073580145),  
          ('suppos', -4.587736246090588),  
          ('lack', -4.520590601432766),  
          ('save', -4.369268990324432),  
          ('minut', -4.25737737396317),  
          ('ridicul', -4.154887819324281),  
          ('annoy', -4.1489427378412955)]
```

```
In [70]: temp = pd.DataFrame(neg_sorted, columns = ["Negative_Words", 'Coef_Magnitude'])
temp['Coef_Magnitude'] = temp['Coef_Magnitude'].abs() # Apply absolute value
fig = px.bar(temp, x="Coef_Magnitude", y="Negative_Words", title='Influential Words in Negative Reviews',
              width=700, height=700, color='Negative_Words')
fig.show()
```

## Influential Words in Negative Reviews



**Words in positive reviews**



```
In [71]: # Zip the names of the features with the features importance
coef_magnitudes_pos = zip(transformer.get_feature_names(), coefs.squeeze()).
```

```
In [72]: # Sort the features in descending order by magnitude
top_100_pos = sorted(coef_magnitudes_pos, key=lambda x: x[1], reverse=True)
top_100_pos
```

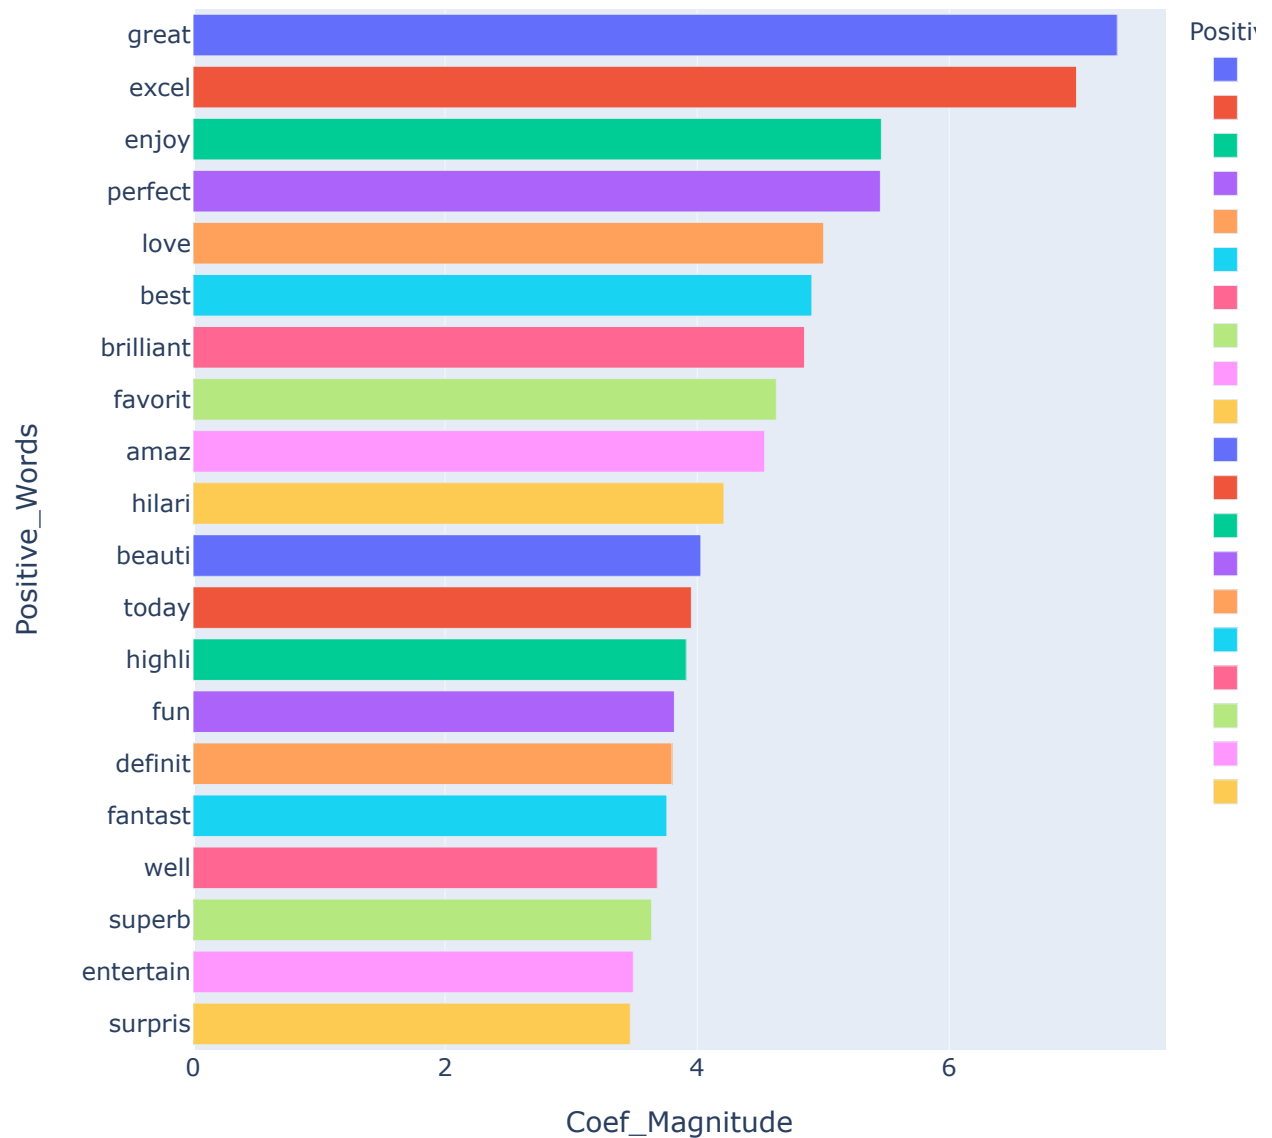
```
Out[72]: [('great', 7.333116055961179),
 ('excel', 7.0107668593828425),
 ('enjoy', 5.461692705850696),
 ('perfect', 5.453807875926427),
 ('love', 5.004253278488765),
 ('best', 4.9104738806299055),
 ('brilliant', 4.852488902344012),
 ('favorit', 4.62837021507443),
 ('amaz', 4.53613492456775),
 ('hilari', 4.212978064617924),
 ('beauti', 4.029569690260818),
 ('today', 3.954248738102151),
 ('highli', 3.913240544320295),
 ('fun', 3.8195655684280467),
 ('definit', 3.8006248798368327),
 ('fantast', 3.759831580093193),
 ('well', 3.683343325715259),
 ('superb', 3.6393337837101436),
 ('entertain', 3.493478317205275),
 ('surpris', 3.469791543086244),
 ('perfectli', 3.391571581526942),
 ('still', 3.3670483766058434),
 ('funniest', 3.329525911624191),
 ('recommend', 3.2406623820935834),
 ('refresh', 3.2083419246120535),
 ('touch', 3.184655907131137),
 ('uniqu', 3.133547532597266),
 ('gem', 3.0735696192589246),
 ('strong', 3.017431167622673),
 ('classic', 2.9872924895192656),
 ('realist', 2.9305204918788785),
 ('awesom', 2.9076940365129396),
 ('simpl', 2.8838150794414887),
 ('especi', 2.8772420956042963),
 ('subtl', 2.8668667522935767),
 ('appreci', 2.7682067093353506),
 ('terrif', 2.722485042419419),
 ('bit', 2.7164610891540124),
 ('good', 2.715784487385702),
 ('greatest', 2.7109445329761157),
 ('alway', 2.6913266144408308),
 ('dvd', 2.656792842075479),
 ('delight', 2.59290663986185),
 ('fascin', 2.5900236426574557),
 ('also', 2.5511719987930284),
 ('thank', 2.546770277237754),
 ('underr', 2.517240014760455),
 ('see', 2.4906795214130297),
 ('masterpiec', 2.4781289092552186),
 ('marvel', 2.4595572614358368),
 ('solid', 2.452166259474853),
 ('wonder', 2.3941688819612827),
 ('human', 2.363460336292522),
```

```
(
    'outstand', 2.3477891531399773),
    ('job', 2.3269732291877188),
    ('world', 2.295604701918964),
    ('atmospher', 2.2719403353736127),
    ('perform', 2.246272594777399),
    ('differ', 2.2292345651803553),
    ('surprisingli', 2.219885415377982),
    ('life', 2.138420142715407),
    ('everyon', 2.126899513103413),
    ('rare', 2.1141027828008863),
    ('memor', 2.1114262183989707),
    ('finest', 2.10269320462038),
    ('brilliantli', 2.096652945797733),
    ('favourit', 2.057850841309165),
    ('intens', 1.9995002200892762),
    ('lot', 1.9834016718375036),
    ('power', 1.9446961329513008),
    ('nevertheless', 1.9409529880737946),
    ('nice', 1.918148352743646),
    ('true', 1.913120013881596),
    ('complaint', 1.9122732880137063),
    ('innoc', 1.9076845648885745),
    ('sweet', 1.8941332616901039),
    ('tear', 1.8904384632132085),
    ('unexpected', 1.8787664840084743),
    ('emot', 1.8569343099612954),
    ('time', 1.8389073829044953),
    ('chanc', 1.837408749734956),
    ('heart', 1.8255005524932897),
    ('noir', 1.8151173261403368),
    ('seen', 1.802153148541798),
    ('keep', 1.778642165417919),
    ('delici', 1.7765334642567352),
    ('tale', 1.7432399670081011),
    ('glad', 1.7359438291401044),
    ('edg', 1.7353695960057385),
    ('stun', 1.7349184092432715),
    ('think', 1.7048208048559648),
    ('chill', 1.6873663991020564),
    ('dream', 1.6831148441257804),
    ('sometim', 1.6803712939956506),
    ('journey', 1.672462859452585),
    ('incred', 1.6639426332867486),
    ('deal', 1.6618218604198307),
    ('extraordinari', 1.6598446237628557),
    ('seat', 1.659237111281219),
    ('pleasantli', 1.6417542553999815)]
```

```
In [73]: coef_magnitudes_pos = zip(transformer.get_feature_names(), coefs.squeeze().
pos_sorted = sorted(coef_magnitudes_pos, key=lambda x: x[1], reverse=True)[
```

```
In [74]: temp = pd.DataFrame(pos_sorted, columns = ["Positive_Words", 'Coef_Magnitude'])
fig = px.bar(temp, x="Coef_Magnitude", y="Positive_Words", title='Influential Words in Positive Reviews',
              width=700, height=700, color='Positive_Words')
fig.show()
```

## Influential Words in Positive Reviews



## Evaluation

Our originally stated hypothesis was:

'A model can be derived to input opinionated language material written about a film or show, so that it can reliably predict positive or negative sentiment.'

After text preprocessing and model selection, the resulting findings are:

1. We can reliably predict movie/show sentiment at a rate of 89.3%. This is significantly better than our model-less baseline of 50%.
2. We have identified an abundance of words with the most significant predictive power from our dataset of 50,000 reviews. With testing, it will be interesting to see if these feature importances are applicable to review material outside of IMDB's database.

## Next Steps

As mentioned above, we may be able to tune for better performance by optimizing the following hyperparameters:

1. Vectorizer 'max features' hyperparameter
2. Inverse of regularization strength 'C' hyperparameter

In [ ]: